



# Fundamentals of Deep Learning

Part 2: How a Neural Network Trains



# Agenda

- Part 1: An Introduction to Deep Learning
- Part 2: How a Neural Network Trains
- Part 3: Convolutional Neural Networks
- Part 4: Data Augmentation and Deployment
- Part 5: Pre-Trained Models
- Part 6: Advanced Architectures

# Recap of the Exercise

What just happened?

Loaded and visualized our data

Edited our data (reshaped, normalized, to categorical)


Created our model

Compiled our model


Trained the model on our data

# Data Preparation

Input as an Array



# An Untrained Model





# A Simpler Model

# A Simpler Model

$$y = mx + b$$

x	y
1	3
2	5





$$m = ?$$
  
$$b = ?$$

# A Simpler Model

$$y = mx + b$$

x	y
1	3
2	5





$$m = ?$$

$$b = ?$$

# A Simpler Model

$$y = mx + b$$

x	y	$\hat{y}$
1	3	4
2	5	3



Start Random

$m = -1$

$b = 5$

## A Simpler Model

$$y = mx + b$$


x	y	$\hat{y}$	$err^2$
1	3	4	1
2	5	3	4

$$MSE =$$

$$2.5$$

$$RMSE =$$

$$1.6$$



$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$


# A Simpler Model

$$y = mx + b$$

x	y	$\hat{y}$	$err^2$
1	3	4	1
2	5	3	4


MSE = 2.5

RMSE = 1.6





```
1  data = [(1, 3), (2, 5)]
2  m = -1
3  b = 5
4
5
6  def get_rmse(data, m, b):
7      """Calculates Mean Square Error"""
8      n = len(data)
9      squared_error = 0
10
11     for x, y in data:
12         # Find predicted y
13         y_hat = m*x+b
14
15         # Square difference between
16         # prediction and true value
17         squared_error += (
18             y - y_hat)**2
19
20     # Get average squared difference
21     mse = squared_error / n
22
23     # Square root for original units
24     return mse ** .5
```



# The Loss Curve





# The Loss Curve



# The Loss Curve




# The Loss Curve




# The Loss Curve

The Gradient	Which direction loss decreases the most
$\lambda$ : The learning rate	How far to travel
Epoch	A model update with the full dataset
Batch	A sample of the full dataset
Step	An update to the weight parameters




# The Loss Curve

The Gradient	Which direction loss decreases the most
$\lambda$ : The learning rate	How far to travel
Epoch	A model update with the full dataset
Batch	A sample of the full dataset
Step	An update to the weight parameters




# Optimizers



- Adam
- Adagrad
- RMSprop
- SGD


# **From Neuron to Network**

# Building a Network




- Scales to more inputs

# Building a Network



- Scales to more inputs
- Can chain neurons

# Building a Network



- Scales to more inputs
- Can chain neurons
- If all regressions are linear, then output will also be a linear regression


# Activation Functions

# Activation Functions

## Linear

$$\hat{y} = wx + b$$


```
1 # Multiply each input  
2 # with a weight (w) and  
3 # add intercept (b)  
4 y_hat = wx+b
```



## ReLU

$$\hat{y} = \begin{cases} wx + b & \text{if } wx + b > 0 \\ 0 & \text{otherwise} \end{cases}$$


```
1 # Only return result  
2 # if total is positive  
3 linear = wx+b  
4 y_hat = linear * (linear > 0)
```



## Sigmoid


$$\hat{y} = \frac{1}{1 + e^{-(wx+b)}}$$

```
1 # Start with line  
2 linear = wx + b  
3 # Warp to - inf to 0  
4 inf_to_zero = np.exp(-1 * linear)  
5 # Squish to -1 to 1  
6 y_hat = 1 / (1 + inf_to_zero)
```




# Activation Functions





Linear




ReLU



Sigmoid



# Activation Functions



# Overfitting



# Overfitting

Why not have a super large neural network?





# Overfitting

Which Trendline is Better?



# Overfitting

Which Trendline is Better?



# Training vs Validation Data

Avoid memorization

## Training data


- Core dataset for the model to learn on

## Validation data

- New data for model to see if it truly understands (can generalize)


## Overfitting

- When model performs well on the training data, but not the validation data (evidence of memorization)
- Ideally the accuracy and loss should be similar between both datasets




# **From Regression to Classification**

# An MNIST Model




# An MNIST Model

During Prediction




# An MNIST Model


During Training





# RMSE For Probabilities?




# RMSE For Probabilities?



# Cross Entropy




# Cross Entropy




$$Loss = -((t(x) \cdot \log(p(x))) + (1 - t(x)) \cdot \log(1 - p(x)))$$

$t(x)$  = target (0 if False, 1 if True)


$p(x)$  = probability prediction of point  $x$



# Cross Entropy



```
1 def cross_entropy(y_hat, y_actual):
2     """Infinite error for misplaced confidence."""
3     loss = log(y_hat) if y_actual else log(1-y_hat)
4     return -1*loss
```








A close-up photograph of a green notebook with horizontal ruling lines. The notebook has a dark green cover visible along the right edge. The pages are white with light blue horizontal lines and vertical red margin lines. The text is positioned in the upper left corner of the slide.

**Bringing it Together**

# The Next Exercise

## The American Sign Language Alphabet





A close-up photograph of a green notebook with horizontal ruling lines. The notebook has a dark green cover visible along the right edge. The pages are white with light blue horizontal lines and vertical red margin lines. The text is positioned in the top-left corner of the first page.

**Let's go!**



# Appendix: Gradient Descent

Helping the Computer Cheat Calculus

## Learning from Error

$$MSE = \frac{1}{n} \sum_{i=1}^n (y - \hat{y})^2 = \frac{1}{n} \sum_{i=1}^n (y - (mx + b))^2$$

$$MSE = \frac{1}{2} ((3 - (m(1) + b))^2 + (5 - (m(2) + b))^2)$$

$$\frac{\partial MSE}{\partial m} = 5m + 3b - 13$$



$$\frac{\partial MSE}{\partial m} = -3$$

$$\frac{\partial MSE}{\partial b} = 3m + 2b - 8$$

$$\frac{\partial MSE}{\partial b} = -1$$

$m = -1$   
 $b = 5$


# The Loss Curve



# The Loss Curve

$$\frac{\partial MSE}{\partial m} = -3$$

$$\frac{\partial MSE}{\partial b} = -1$$




# The Loss Curve

$$\frac{\partial MSE}{\partial m} = -3$$

$$\frac{\partial MSE}{\partial b} = -1$$

$$m := m - \lambda \frac{\partial MSE}{\partial m}$$

$$b := b - \lambda \frac{\partial MSE}{\partial b}$$



# The Loss Curve


$$\frac{\partial MSE}{\partial m} = -3$$

$$\frac{\partial MSE}{\partial b} = -1$$

$$m := m - \lambda \frac{\partial MSE}{\partial m}$$

$$\lambda = 2$$

$$b := b - \lambda \frac{\partial MSE}{\partial b}$$



# The Loss Curve


$$\frac{\partial MSE}{\partial m} = -3$$

$$\frac{\partial MSE}{\partial b} = -1$$

$$m := m - \lambda \frac{\partial MSE}{\partial m}$$

$$b := b - \lambda \frac{\partial MSE}{\partial b}$$

$$\lambda = .1$$



# The Loss Curve

$$\lambda = .1$$

$$m := -1 + 3\lambda = -0.7$$

$$b := 5 + \lambda = 5.1$$

